MUST KNOW

RAG

Interview Questions

12	69	∞
SECTIONS	QUESTIONS	VALUE

Chandra Sekhar

AI Engineer & Trainer

Bundle link in Caption - Enroll Now

WHAT'S INSIDE

12 Topics · 69 Questions

RAG Fundamentals & Core Concepts
6 questions

RAG Architecture & Workflow
5 questions

Document Loading & Data Ingestion
7 questions

Text Splitting & Chunking Strategies
9 questions

Embeddings & Vector Stores
5 questions

Vector Search & Retrieval Techniques
8 questions

Vector Store Operations & Debugging
3 questions

Prompting & Context Injection
5 questions

Evaluation, Monitoring & Observability
6 questions

Performance, Cost & Scalability
5 questions

Security, Privacy & Enterprise Readiness
5 questions

Advanced & Agentic RAG (Bonus)
5 questions

Every question here has been asked in real AI Engineer interviews.

SECTION 1

RAG Fundamentals & Core Concepts

1. What is Retrieval Augmented Generation (RAG), and why is it critical for LLMs?
2. What are the limitations of a pure LLM system without RAG?
3. How is RAG different from fine-tuning an LLM?
4. When should you prefer RAG over fine-tuning?
5. What types of problems are best solved using RAG?
6. What kinds of data sources can be used in a RAG system?

Bundle link in Caption - Enroll Now

SECTION 2

RAG Architecture & Workflow

7. What are the two main components of a RAG application architecture?

8. How does the indexing process work, and why is it essential?

9. What is the step-by-step workflow of a complete RAG pipeline?

10. What is the specific role of the Retrieval component?

11. What is the goal of integrating retrieval with generation?

Bundle link in Caption - Enroll Now

SECTION 3

Document Loading & Data Ingestion

12 What is the primary purpose of DocumentLoaders?

13 How does WebBaseLoader function within the LangChain ecosystem?

14 What are the roles of WebBaseLoader and RecursiveCharacterTextSplitter?

15 How can you customize the HTML parsing process in WebBaseLoader?

16 What is the specific role of BeautifulSoup in document ingestion?

17 What is a SoupStrainer and how does it improve efficiency?

18 Why is customized HTML parsing critical for a RAG system?

Bundle link in Caption - Enroll Now

SECTION 4

Text Splitting & Chunking Strategies

19 Why is it necessary to split documents into smaller chunks?

20 Why must long documents be split before embedding?

21 What is chunk overlap and why is it used?

22 How does chunk size impact retrieval accuracy?

23 What trade-offs exist between larger chunks and smaller chunks?

24 How does poor chunking lead to hallucinations?

25 Why is RecursiveCharacterTextSplitter recommended for generic text?

26 How does the `add_start_index=True` parameter

Bundle link in Caption - Enroll Now

SECTION 5

Embeddings & Vector Stores

28 What is the purpose of embedding text and storing it in a vector store?

29 How does the choice of embedding model affect RAG performance?

30 What happens if different embedding models are used for indexing and querying?

31 What are the key components needed to build and query a vector store?

32 What is Cosine Similarity and why is it used in vector searches?

Bundle link in Caption - Enroll Now

SECTION 6

Vector Search & Retrieval Techniques

33 What is the specific role of a Retriever in LangChain?

34 How does the VectorStoreRetriever function?

35 What does the `k` parameter control in vector retrieval?

36 What is the difference between similarity search and MMR (Max Marginal Relevance)?

37 What is hybrid search and why is it useful in RAG systems?

38 How do metadata filters improve retrieval quality?

39 What is reranking and why is it important in advanced RAG pipelines?

Bundle link in Caption - Enroll Now

SECTION 7

Vector Store Operations & Debugging

41 What are common challenges when working with the Chroma vector store?

42 How do you inspect the contents of a vector store?

43 What is the step-by-step workflow to embed and store documents?

Bundle link in Caption - Enroll Now

SECTION 8

Prompting & Context Injection

44 How is retrieved context injected into the prompt?

45 What is context stuffing and why is it dangerous?

46 How do you design prompts to prevent hallucinations in RAG?

47 What happens if retrieved documents contradict each other?

48 How do you limit the LLM to answer only from retrieved context?

Bundle link in Caption - Enroll Now

SECTION 9

Evaluation, Monitoring & Observability

49 How do you evaluate the quality of a RAG system?

50 What metrics are used to measure retrieval quality?

51 How do you detect hallucinations in a RAG pipeline?

52 Why is configuring LangSmith crucial for complex applications?

53 How does LangSmith help in debugging RAG applications?

54 What are common failure modes in RAG systems?

Bundle link in Caption - Enroll Now

SECTION 10

Performance, Cost & Scalability

55 How do you reduce latency in a RAG system?

56 How do you optimize embedding costs?

57 What caching strategies are used in RAG pipelines?

58 How do you handle large-scale document updates?

59 How do you design a RAG system for production use?

Bundle link in Caption - Enroll Now

SECTION 11

Security, Privacy & Enterprise Readiness

60 How does RAG help with private or sensitive data?

61 What are the security risks in a RAG-based system?

62 How do you prevent data leakage in RAG responses?

63 How do access controls work in enterprise RAG systems?

64 How do you ensure compliance when using RAG?

Bundle link in Caption - Enroll Now

SECTION 12

Advanced & Agentic RAG (Bonus)

65 What is Agentic RAG?

66 How does RAG work with tools and agents?

67 What is self-reflection or self-correction in RAG systems?

68 How does multi-step retrieval improve answer quality?

69 What is RAG with feedback loops?

Bundle link in Caption - Enroll Now

Want Access to the Entire AI Interview Bundle?

13	700+	Interview
COURSES	QUESTIONS	READY

🎯 INTRODUCING

The AI Interview Master Bundle

13 Interview Courses 650+ Questions

- ✓ **13 Structured Interview Courses**
  Topic-wise, end-to-end prep
- ✓ **650+ Carefully Curated Questions**
  No fluff — interview-ready
- ✓ **Beginner → Advanced**
  One structured path

Check Next slide for more details

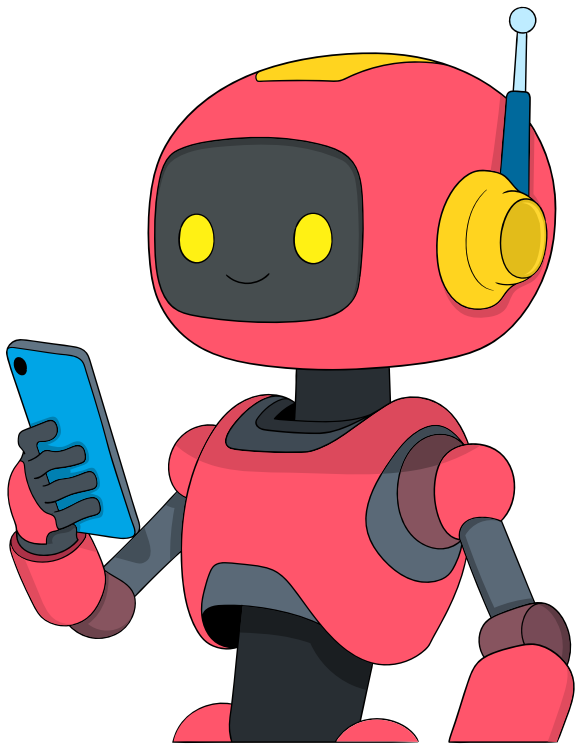

AI INTERVIEW MASTER BUNDLE

Overview of AI Interview Mastery Bundle

13 modules • 675+ questions

1. **Python For GenAI** — 55 questions
2. **Statistics & Math for ML** — 30 questions
3. **Machine Learning Fundam...** — 70 questions
4. **Deep Learning Basics** — 25 questions
5. **LLM Core Concepts** — 50 questions
6. **RAG Interview Mastery** — 69 questions
7. **LangChain Core Concepts** — 134 questions
8. **LangGraph Core Concepts** — 49 questions
9. **Agentic AI** — 50 questions
10. **CrewAI** — 53 questions
11. **AutoGen** — 28 questions
12. **FastAPI (Backend for AI)** — 32 questions
13. **Docker for AI Engineers** — 30 questions

• **Access bundle link from caption**

GET HIRED

The only bundle you'll ever need.

Stop scattered prep. Start structured learning.

13 modules · 645 questions · built for AI engineers

🔗 **Link of Bundle in caption**

Chandra Sekhar

AI Engineer and Trainer

Follow for daily AI interview questions, tips & career advice.

Follow on LinkedIn →